AI Frontier Model Providers & Releases Matrix

A comprehensive mapping of models released in the past 12 months with token metrics, context windows, and target use cases.

Model Provider	Model	Release Date	Cut-off Date	Cost per 1M Tokens (In / Out)	Context Window	Type	When to Use / Core Use Case
Anthropic	Claude Opus 5.5	Sep 22, 2026	June 2026	$4.00 / $20.00	1,000,000	Closed-source Reasoning	Complex, long-running agentic coding tasks and enterprise-grade knowledge curation deliverables.
OpenAI	GPT-6 Sol	Sep 22, 2026	Apr 20, 2026	$2.00 / $10.00	1,050,000	Closed-source Multimodal Reasoning	High-end professional everyday automation, complex computer control loops, and software engineering planning.
Google	Gemini 3.5 Flash	May 19, 2026	Jan 2025	$1.50 / $9.00	1,048,576	Closed-source Multimodal	High-efficiency native audio, video, and PDF processing requiring rapid, multi-tier parallel agent execution.
xAI	Grok 4.7	Sep 21, 2026	May 2026	$2.00 / $6.00	500,000	Closed-source Vision Reasoner	Deep text and visual problem-solving demanding real-time X platform search integration and high-speed execution.
DeepSeek	DeepSeek-V4.1-Flash	Sep 10, 2026	—	$0.30 / $1.20	1,000,000	Open-API Mixture-of-Experts	Maximum cost efficiency for massive scale tool-calling environments, codebase ingestion, and high-concurrency loops.
MiniMax	MiniMax M3.1 Flash	Oct 19, 2026	—	$0.15 / $0.60	1,000,000	Open-weights Multimodal Reasoner	Long-horizon video/image inputs paired with structural multi-turn interaction loops without long-context latency penalties.
Alibaba Cloud	Qwen3.5-Plus	Feb 16, 2026	Feb 2026	$0.40 / $2.40	1,000,000	Proprietary Multimodal API	Advanced text/image/video multi-step logic, complex agentic programming, and localized long-form text synthesis.

Model Provider	Model	Release Date	Cut-off Date	Cost per 1M Tokens (In / Out)	Context Window	Type	When to Use / Core Use Case
Perplexity	Decider v1 27B	Oct 2026	—	$0.02 / $0.00	—	Closed-source Router	Intended strictly for high-throughput decision gates, fast fallback model routing, and initial pipeline query filtering.
Mistral AI	Mistral Large 4	Late 2026	—	$0.68 / $2.09	1,000,000	Proprietary Large Context	Enterprise data ingestion, large archive analytics, and robust multilingual translation or long document summarization.
Z.AI	GLM-5.1	Late 2026	—	$0.60 / $1.60	200,000	Closed-source Conversational	Fast localized enterprise workflows, rapid function calling parameters, and multi-turn workflow scheduling.